

```
In [1]: import pandas as pd
df = pd.read_csv("C:\\Users\\Sushmitha T\\Downloads\\Advertising Dataset.csv")
```

```
In [8]: feature_cols=['TV','Radio','Newspaper']
```

```
In [10]: x=df[feature_cols]
x=df[['TV','Radio','Newspaper']]
x.head()
```

```
Out[10]:
```

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4

```
In [11]: print(type(x))
print(x.shape)

<class 'pandas.core.frame.DataFrame'>
(200, 3)
```

```
In [13]: y=df['Sales']
y=df.Sales
y.head()
```

```
Out[13]:
```

0	22.1
1	10.4
2	9.3
3	18.5
4	12.9

Name: Sales, dtype: float64

```
In [14]: print(type(y))
print(y.shape)

<class 'pandas.core.series.Series'>
(200,)
```

Splitting x and y training and testing sets

```
In [22]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=1)
```

```
In [23]: print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(150, 3)
(50, 3)
(150,)
(50,)
```

Linear Regression in scikit-learn

```
In [24]: from sklearn.linear_model import LinearRegression
```

```
In [26]: linreg=LinearRegression()
linreg.fit(x_train,y_train)
```

```
Out[26]:
```

LinearRegression ⓘ ?

LinearRegression()

Interpreting model coefficients

```
In [27]: print(linreg.intercept_)
print(linreg.coef_)

2.8769666223179335
[0.04656457 0.17915812 0.00345046]
```

```
In [28]: list(zip(feature_cols,linreg.coef_))
```

```
Out[28]: [('TV', np.float64(0.046564567874150295)),
          ('Radio', np.float64(0.1791581224508883)),
          ('Newspaper', np.float64(0.003450464711180402))]
```

```
In [29]: y_pred=linreg.predict(x_test)
```

```
In [30]: true=[100,50,30,20]
pred=[90,50,50,30]
```

```
In [17]: from sklearn import metrics
true = [100,50,30,20]
pred = [90,50,50,30]
print((10+0+20+10)/4.)
print(metrics.mean_absolute_error(true, pred))
```

```
10.0
10.0
```

```
In [18]: print((10**2 + 0**2 + 20**2 + 10**2)/4.)
from sklearn import metrics
print(metrics.mean_squared_error(true,pred))
```

```
150.0
150.0
```

central tendency

```
In [19]: mean_values = df.mean()
print("Mean:\n", mean_values)
```

```
Mean:
Unnamed: 0    100.5000
TV           147.0425
Radio        23.2640
Newspaper    30.5540
Sales        14.0225
dtype: float64
```

```
In [20]: median_values = df.median()
print("Median:\n", median_values)
```

```
Median:
Unnamed: 0    100.50
TV           149.75
Radio         22.90
Newspaper     25.75
Sales         12.90
dtype: float64
```

```
In [21]: mode_values = df.mode()
print("Mode:\n", mode_values)
```

```
Mode:
   Unnamed: 0    TV  Radio  Newspaper  Sales
0           1  17.2   4.1         8.7    9.7
1           2  76.4   5.7         9.3   NaN
2           3 109.8   NaN        25.6   NaN
3           4 177.0   NaN         NaN   NaN
4           5 184.9   NaN         NaN   NaN
..         ...   ...   ...         ...   ...
195        196   NaN   NaN         NaN   NaN
196        197   NaN   NaN         NaN   NaN
197        198   NaN   NaN         NaN   NaN
198        199   NaN   NaN         NaN   NaN
199        200   NaN   NaN         NaN   NaN
```

[200 rows x 5 columns]

Measures of Dispersion

```
In [22]: range_values = df.max() - df.min()
print("Range:\n", range_values)
```

```
Range:
Unnamed: 0    199.0
TV           295.7
Radio         49.6
Newspaper    113.7
Sales        25.4
dtype: float64
```

```
In [23]: variance_values = df.var()
print("Variance:\n", variance_values)
```

```
Variance:
Unnamed: 0    3350.000000
TV           7370.949893
Radio        220.427743
Newspaper    474.308326
Sales        27.221853
dtype: float64
```

```
In [24]: std_dev_values = df.std()
print("Standard Deviation:\n", std_dev_values)
```

```
Standard Deviation:
Unnamed: 0    57.879185
TV           85.854236
Radio        14.846809
Newspaper    21.778621
Sales        5.217457
dtype: float64
```

```
In [ ]:
```

Identifying the intercept (c) and coefficients (m_1, m_2, m_3, \dots):

```
In [23]: X=df.drop('Sales',axis=1)
y=df['Sales']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
intercept = model.intercept_
coefficients = model.coef_

# Print the intercept and coefficients
print(f"Intercept (c): {intercept}")
print(f"Coefficients (m1, m2, m3, ...): {coefficients}")
```

```
Intercept (c): 2.906527086361816
Coefficients (m1, m2, m3, ...): [0.00064359 0.04471835 0.18925118 0.00304577]
```

```
In [ ]:
```

```
In [25]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
data = pd.read_csv('C:\\Users\\Sushmitha T\\Downloads\\Advertising Dataset.csv')
X = data.drop('Sales', axis=1)
Y = data['Sales']
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.2,random_state=42)
from sklearn.linear_model import LinearRegression
linreg = LinearRegression()
linreg.fit(X_train, Y_train)

Y_pred = linreg.predict(X_test)

mse = mean_squared_error(Y_test, Y_pred)
r2 = r2_score(Y_test, Y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared (R²) score: {r2}")

comparison = pd.DataFrame({'Actual': Y_test, 'Predicted': Y_pred})
print(comparison.head())
```

```
Mean Squared Error: 3.199004468588908
R-squared (R²) score: 0.8986489151417079
   Actual  Predicted
95    16.9   16.412277
15    22.4   20.843193
30    21.4   21.511869
158    7.3   10.653100
128   24.7   22.124058
```

```
In [27]: import numpy as np
from sklearn.metrics import mean_squared_error
rmse = np.sqrt(mean_squared_error(Y_test, Y_pred))
print(f"Root Mean Squared Error (RMSE): {rmse}")
print(f"R-squared (R²) score: {r2}")
print("\n")
```

```
Root Mean Squared Error (RMSE): 1.7885761008659677
R-squared (R²) score: 0.8986489151417079
```

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js